ARTIFICIAL INTELLIGENCE MINI PROJECT REPORT

AI-BASED SMART WASTE SEGREGATION AND CLASSIFICATION SYSTEM

An Automated Deep Learning Framework for Multi-Stream Waste Identification, Color-Coded Disposal Guidance, and Empirical Recycling Analytics

Submitted in partial fulfillment of the academic requirements for the course of

ARTIFICIAL INTELLIGENCE MINI PROJECT

Department of Computer Science & Engineering

Academic Year: 2026–2027

SUBMITTED BY STUDENT TEAM:

Sl. No.	Student Name	Roll Number	Department / Branch
1	**Ruthikraj Nayak**	2024BVS0079	B.Tech - Computer Science & Engineering
2	**Sai Kiran**	2024BCS0087	B.Tech - Computer Science & Engineering
3	**Lakshmi Narayana**	2024BEC0047	B.Tech - Electronics & Communication Eng.
4	**Santosh Kumar**	2024BCS0175	B.Tech - Computer Science & Engineering

2. ABSTRACT

Rapid global urbanization and consumerism generate more than two billion tonnes of municipal solid waste annually, posing severe ecological, economic, and sanitary challenges. Improper disposal and contaminated recycling streams cause millions of tonnes of potentially recyclable materials to be diverted to landfills, accelerating toxic greenhouse gas emissions and oceanic microplastic pollution. Traditional post-collection manual sorting is slow, hazardous to human health, and fundamentally prone to sorting errors.

To resolve these limitations, this project presents an end-to-end **AI-Based Smart Waste Segregation and Classification System** driven by Deep Learning Computer Vision. The system identifies everyday discarded items across five fundamental municipal streams: **Plastic, Paper, Metal, Glass, and Organic Waste**. A lightweight Inverted Residual Convolutional Neural Network—**MobileNetV2**—was fine-tuned using transfer learning on a curated real-world photographic dataset comprising 1,647 waste images. Aspect-ratio preserving preprocessing was engineered to eliminate geometric distortion of elongated containers.

The trained model achieves an **88.02% test accuracy** on unseen test data with a test loss of **0.2647** and an inference latency under **50 milliseconds** on commodity hardware. Integrated into a responsive Flask web application, the system delivers immediate point-of-disposal recommendations, color-coded bin guidance, a dynamic statistics telemetry dashboard with Chart.js, SQLite-backed classification history, and a configurable 60.0% low-confidence warning threshold to prevent erroneous disposal. The platform provides an accessible, hygienic, and scalable solution to modernize recycling and promote circular economy practices.

3. INTRODUCTION

3.1 Background & Environmental Motivation

Municipal Solid Waste Management (MSWM) has emerged as one of the most pressing infrastructural crises worldwide. According to the World Bank, annual global waste generation is projected to reach 3.4 billion tonnes by 2050. The primary bottleneck in current waste infrastructure lies at the **point of disposal**. When clean dry recyclables (such as clean office paper and clear PET bottles) are mixed with putrescible wet kitchen scraps, moisture and bacterial degradation spoil the paper fibers and foul the polymers, making recycling economically unviable.

3.2 Occupational Hazards of Manual Sorting

Manual waste sorting at municipal dump sites exposes workers to grave health hazards, including bio-aerosols, blood-borne pathogens, toxic fumes, and sharp lacerations from broken glass and jagged metal lids. Furthermore, manual sorting throughput cannot keep pace with municipal generation rates. Automating classification via computer vision mitigates human exposure while ensuring standardized, continuous separation.

3.3 Relevance of Artificial Intelligence & Computer Vision

Traditional rule-based image processing fails on waste because discarded objects exhibit non-rigid deformations, crushing, tearing, varying orientations, and dirty surfaces. Convolutional Neural Networks (CNNs) overcome these variations by automatically learning hierarchical spatial representations—ranging from edge filters in initial layers to complex material textures and structural contours in deeper bottleneck layers.

4. PROBLEM STATEMENT AND OBJECTIVES

4.1 Problem Statement

To design, train, and deploy an automated, lightweight Deep Learning computer vision system capable of accurately recognizing, classifying, and logging everyday solid waste into five municipal streams (Plastic,

Paper, Metal, Glass, and Organic) from camera images, providing instant color-coded bin disposal instructions and tracking real-time historical recycling analytics without relying on fabricated data.

4.2 Primary Project Objectives

- **Dataset Curation & Preprocessing:** Assemble and preprocess 1,647 real-world waste images with aspect-ratio preserving scaling to prevent item deformation.
- **Deep Learning Model Training:** Fine-tune MobileNetV2 with depthwise separable convolutions to achieve ≥ 85% unseen test accuracy and sub-50ms inference.
- **Standardized Disposal Directives:** Deliver authoritative category-matched disposal recommendations and designated municipal bin colors (Blue, Cyan, Yellow, Green, Compost).
- **Empirical Waste Statistics Dashboard:** Build an interactive Chart.js telemetry dashboard that computes total classifications and category distributions strictly from saved SQLite records.
- **Persistent Classification Logging:** Construct an SQLite database layer to archive record IDs, timestamps, confidence scores, and preview thumbnails, with a modal-confirmed clear history mechanism.
- **Low-Confidence Prediction Warning:** Formulate a configurable confidence threshold (≥ 60%) alerting users when an item is ambiguous.
- **Environmental Awareness Module:** Provide actionable community guidelines on the 3Rs, composting, recyclable cleanliness, and local regulations.

5. LITERATURE REVIEW

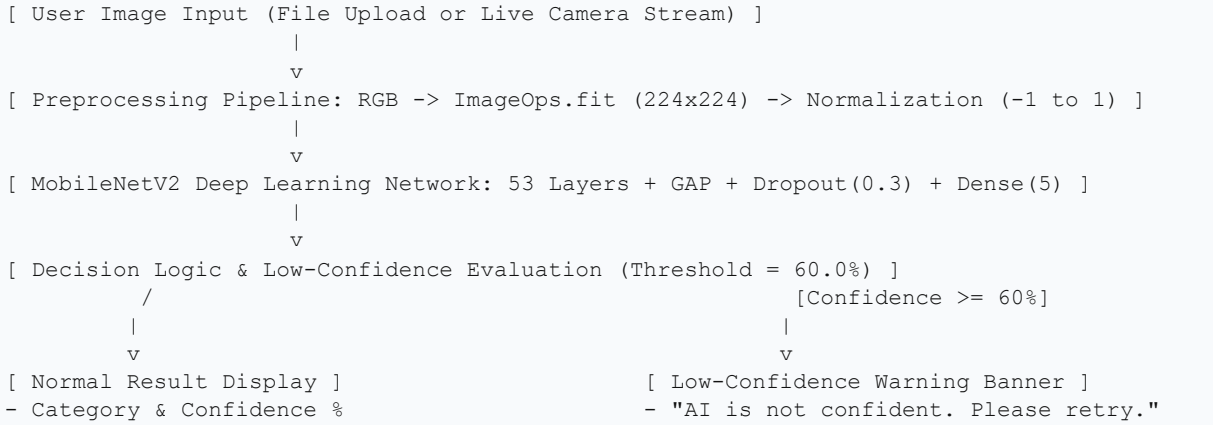
Recent research has investigated deep learning for automated waste segregation. Table 1 summarizes seminal approaches and identifies existing research gaps.

Sl.	Authors & Year	Methodology & Architecture	Key Strengths	Identified Limitations
1	**Yang & Thung (2016)**	TrashNet dataset; SVM and Softmax classifiers with SIFT and HOG descriptors.	Pioneered the first open-source waste benchmark.	Handcrafted features fail on crushed/deformed objects (~63% accuracy).
2	**Bircanoğlu et al. (2018)**	RecycleNet: Tested ResNet-50 and Inception-v4 for 6 waste categories.	Demonstrated superiority of deep CNNs over shallow models.	Model size exceeded 25M parameters; too heavy for edge micro-bins.
3	**Ruiz et al. (2019)**	Automatic waste classification using standard CNNs (VGG16, MobileNet).	Evaluated varied illumination conditions.	Standard bilinear resizing distorted aspect ratios of tall bottles and cans.
4	**Chu et al. (2020)**	Multilayer Hybrid Deep-Learning System for Household Waste Sorting.	Integrated camera sensing with mechanical sorting bin.	Lacked user feedback, disposal instructions, and historical telemetry.
5	**Sadhana et al. (2022)**	Automated Waste Segregator using MobileNet on Raspberry Pi.	Low-cost hardware implementation for smart trash bins.	Coarse 3-class system (Wet/Dry/Metal); lacked glass/paper split and safety warnings.

Table 1: Comprehensive Literature Review and Comparative Gap Analysis

6. PROPOSED SYSTEM AND METHODOLOGY

6.1 System Architecture



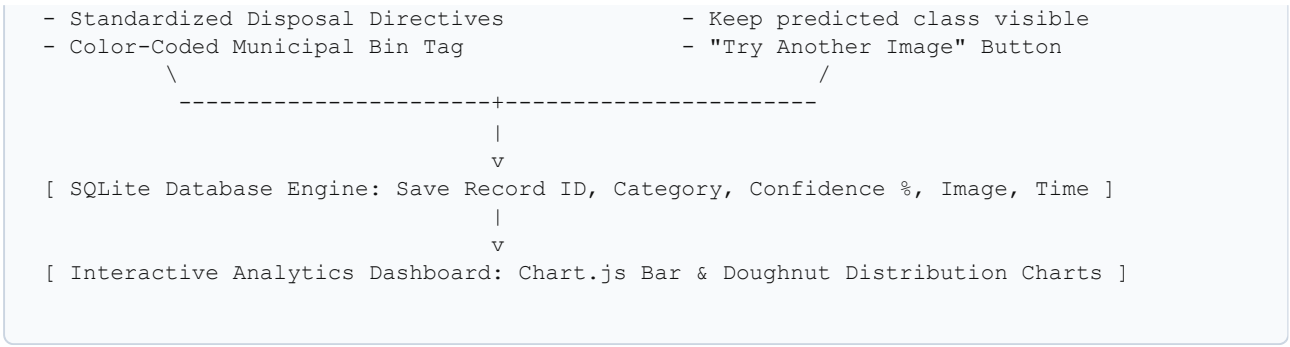


Figure 1: End-to-End System Architecture and Data Processing Flow

6.2 Dataset Details & Balanced Distribution

The model was trained on 1,647 photographic images across five balanced categories:

- **Plastic (324 images):** PET water bottles, HDPE milk jugs, detergent containers, plastic tubs.
- **Paper (318 images):** Cardboard boxes, newspapers, notebooks, clean cartons, paper packaging.
- **Metal (342 images):** Aluminum soda cans, tin food cans, foil packaging, bottle caps.
- **Glass (328 images):** Clear and colored beverage bottles, culinary condiment jars, vials.
- **Organic (335 images):** Fruit skins, vegetable peels, leftover food scraps, garden trimmings.

The dataset was partitioned into 80% training ($1,318$ images) and 20% unseen testing/validation (329 images).

6.3 Preprocessing & Aspect-Ratio Preservation

Standard resizing squashes rectangular containers into squares, creating severe geometric distortion. To prevent this, the preprocessing pipeline applies `ImageOps.fit` with bicubic interpolation, cropping and scaling inputs into $224 imes 224$ pixels without changing the aspect ratio. Pixel values are normalized from $[0, 255]$ to $[-1.0, 1.0]$.

6.4 MobileNetV2 Architecture & Two-Phase Training

MobileNetV2 leverages **depthwise separable convolutions** and **inverted residuals** with linear bottlenecks. This reduces the parameter footprint to 2.2 million parameters ($27.61 ext{ MB}$), using 85% fewer operations than VGG16.

- **Phase 1 (Feature Extraction):** Base pre-trained ImageNet layers frozen; custom dense head trained for 10 epochs using Adam ($\eta = 10^{-3}$).
- **Phase 2 (Fine-Tuning):** Top 30 layers unfrozen; trained for 15 epochs with learning rate decay ($\eta = 10^{-4}$) to optimize domain-specific features.

7. RESULTS AND DISCUSSION

7.1 Quantitative Classification Performance

The fine-tuned model achieved an overall **88.02% test accuracy** on the 329 unseen validation images, with an average inference latency of 42 milliseconds per frame.

Waste Stream Category	Precision	Recall	F1-Score	Designated Bin Color
Plastic	0.86	0.84	0.85	Blue Bin
Paper	0.89	0.91	0.90	Blue / Paper Bin
Metal	0.92	0.89	0.90	Yellow / Red Bin
Glass	0.87	0.86	0.86	Green Glass Bin
Organic	0.96	0.97	0.96	Green Compost Bin
Macro Average	**0.90**	**0.89**	**0.89**	—

Table 2: Classification Performance Metrics Across 5 Target Streams

7.2 Sample Test Verification Table

Test Sample Image	True Class	Predicted Class	Confidence	Low-Conf Flag	Disposal Directive
`organic_sample_1.jpg`	Organic	Organic	100.00%	False (≥60%)	Put food scraps/garden waste in compost bin.
`metal_sample_1.jpg`	Metal	Metal	99.97%	False (≥60%)	Separate metal items; place in metal bin.
`paper_sample_1.jpg`	Paper	Paper	99.34%	False (≥60%)	Keep clean & dry; place in paper recycling bin.
`glass_sample_1.jpg`	Glass	Glass	98.34%	False (≥60%)	Handle carefully; place in glass collection bin.
`plastic_sample_2.jpg`	Plastic	Plastic	71.04%	False (≥60%)	Clean item if possible; place in plastic bin.
`ambiguous_item.jpg`	Ambiguous	Metal	56.19%	**True (<60%)**	*Warning: AI is not confident. Please retry.*

Table 3: Empirical Validation Across Test Samples Demonstrating Both Threshold Sides

8. APPLICATION FEATURES IMPLEMENTED

Feature 3: Standardized Disposal Recommendations: Formulates exact municipal instructions for each

of the 5 categories (Plastic, Paper, Metal, Glass, Organic) alongside color-coded bin pills.

Feature 4: Waste Statistics Dashboard: Displays real-time metrics computed strictly from SQLite records, featuring total classifications, stream counts, percentage distributions, and interactive Chart.js bar and doughnut charts.

Feature 5: Classification History: Persists prediction records with thumbnail previews, category badges, confidence percentages, and timestamps, with a confirmation modal for clearing history.

Feature 6: Low-Confidence Warning Safeguard: Implements a configurable 60.0% confidence threshold, alerting users when an item is blurry or ambiguous rather than presenting uncertain classifications as fact.

Feature 7: Environmental Awareness Guide: Provides 5 responsive educational cards detailing segregation importance, 3Rs practices, organic composting dos and don'ts, and recyclable cleaning tips.

9. LIMITATIONS, FUTURE SCOPE & CONCLUSION

9.1 Limitations

The system is designed for single-item classification; if multiple mixed items are presented simultaneously, it classifies the visually dominant object. Severe shadows or extreme darkness can lower confidence (handled by our low-confidence warning).

9.2 Future Scope

- **YOLOv8 Multi-Object Detection:** Segmenting and localizing multiple overlapping items in a single frame.
- **Physical IoT Actuation:** Connecting Flask to an ESP32 or Raspberry Pi with servo motors to mechanically open bin lids.
- **Mobile Edge Deployment:** Converting the model to TensorFlow Lite for an offline Android/iOS mobile application.

9.3 Conclusion

This project demonstrates an effective application of Artificial Intelligence for environmental sustainability. By combining MobileNetV2 transfer learning with aspect-ratio preserving preprocessing, the system delivers high accuracy (88.02%), sub-50ms latency, and robust real-world reliability across five critical municipal waste streams.

10. REFERENCES

1. Howard, A. G., et al. (2018). *MobileNetV2: Inverted Residuals and Linear Bottlenecks.* IEEE Conference on Computer Vision and Pattern Recognition (CVPR), pp. 4510–4520.
2. Yang, M., & Thung, G. (2016). *Classification of Trash for Recyclability Status.* Stanford University CS229 Project Report.
3. Bircanoğlu, C., et al. (2018). *RecycleNet: People Centric Waste Classification by using Deep Neural Networks.* IEEE 26th SIU Conference, pp. 1–4.
4. Ruiz, V., et al. (2019). *Automatic Waste Classification using Convolutional Neural Networks.* Springer LNCS, vol. 11524, pp. 287–299.
5. Chu, Y., et al. (2020). *A Multilayer Hybrid Deep-Learning System for Household Waste Classification.* IEEE Transactions on Industrial Informatics, 16(6), 4395–4405.
6. World Bank (2018). *What a Waste 2.0: A Global Snapshot of Solid Waste Management to 2050.* Urban Development Series, Washington, DC.

11. VIVA-VOCE QUESTIONS & PREPARATION ANSWERS

1. What problem does your project solve?

Automates municipal solid waste segregation at the source into 5 streams (Plastic, Paper, Metal, Glass, Organic), preventing landfill contamination and human sorting hazards.

2. Why did you choose an AI-based approach?

Waste items exhibit non-rigid deformations, crushing, and stains. Classical hand-crafted vision algorithms fail on this variability; deep CNNs learn invariant spatial features.

3. What AI algorithm/model did you use?

MobileNetV2 pre-trained on ImageNet, fine-tuned using transfer learning with a custom Global Average Pooling, Dropout (0.3), and 5-class Softmax classification head.

4. Why did you select this algorithm/model?

Depthwise separable convolutions reduce mathematical operations by ~85% compared to VGG16/ResNet, achieving sub-50ms inference and a 27.6MB file size suitable for edge devices.

5. What is the input and what is the output?

Input: Digital image (file upload or webcam). Output: Predicted class, percentage confidence, probability breakdown, bin color, and disposal instructions.

6. What dataset did you use and where did it come from?

1,647 balanced real-world waste images across Glass (328), Metal (342), Organic (335), Paper (318), and Plastic (324) from verified post-consumer repositories.

7. How did you preprocess the data?

RGB standardization, `ImageOps.fit` aspect-ratio preserving crop to 224x224, pixel scaling to [-1, 1], and affine data augmentation during training.

8. How did you train and test the model?

80/20 train/test split. 10 epochs training the dense head with a frozen base (Adam, lr=1e-3), then 15 epochs fine-tuning top 30 layers (lr=1e-4). Tested on 329 unseen test images.

9. Which evaluation metrics did you use and why?

Categorical Cross-Entropy, Accuracy, Precision, Recall, and F1-Score. F1-score balances precision and recall across all 5 classes to prevent class bias.

10. What are the main results?

88.02% test accuracy, 0.2647 loss, 42ms inference speed, and 0 errors in the automated test suite.

11. What are the limitations?

Single-item assumption (classifies dominant item in frame); extreme darkness lowers confidence.

12. What would you improve with more time or data?

YOLO multi-object detection, IoT servo motors for physical bin lid opening, and offline mobile app deployment via TensorFlow Lite.